

VELLORE INSTITUTE OF TECHNOLOGY

School of Computer Science and Engineering (SCOPE) | Fall 2026-2027

Course: MDI3003 Advanced Predictive Analytics
Faculty: Dr. Durgesh Kumar, Assistant Professor (Senior)
Student Reg: 23MID0344
NAME: AVINASH A

Experiment 07: Constructing a Recommendation System from Customer Transaction Data using Random Forest

A Chronologically Leakage-Free Supervised Candidate-Scoring, Top-K Ranking, and Advanced Collaborative Benchmark Report

EXECUTIVE ABSTRACT

This laboratory investigation formulates customer product recommendation as a supervised pair-wise candidate-scoring and ranking problem over implicit-feedback retail transactions. Using the 541,909-transaction UCI Online Retail benchmark, we enforce strict chronological splitting (50%, 70%, 85% quantiles) to eliminate temporal data leakage. A bounded candidate universe of N=1,000 high-velocity products achieves a Candidate Recall of 63.76% on locked future test customers. Leakage-safe feature engineering derives 20 customer RFM, item velocity, and historical co-purchase interaction signals. A tuned Random Forest classifier achieves a test NDCG@10 of **0.3173** and Recall@10 of **0.1716** (Hit Rate@10: **82.53%**), substantially outperforming both the global Popularity Baseline (NDCG@10: 0.0828, Recall@10: 0.0317) and an Item-Item Collaborative Filtering benchmark (NDCG@10: 0.1503). User-level bootstrap 95% confidence intervals confirm statistical significance. An in-depth five-case audit, feature ablation, catalog coverage analysis, and discussion of ethical/privacy boundaries establish production readiness.

DOCUMENT TABLE OF CONTENTS

- | | |
|--|--|
| 1. Problem Statement & Business Context | 9. Top-K Recommendation & Evaluation Suite |
| 2. Laboratory Positioning & Predictive Translation | 10. Core Empirical Results & Model Comparison |
| 3. Data Schema, Provenance & Cleaning Integrity | 11. Advanced CF Benchmark & Uncertainty Quantification |
| 4. Chronological Splitting & Leakage Prevention | 12. Feature Ablation & Sensitivity Studies (E4, E5) |
| 5. Candidate Universe & Candidate-Recall Audit | 13. Five-Case Deep-Dive Diagnostic Audit |
| 6. Leakage-Safe Feature Engineering Architecture | 14. Error, Cold-Start & Popularity Bias Analysis |
| 7. Negative Sampling & Supervised Formulation | 15. System Efficiency & Serving Benchmarks |
| 8. Model Training & Hyperparameter Tuning | 16. Comprehensive Discussion & Viva Key Points |

1. PROBLEM STATEMENT AND BUSINESS CONTEXT

Modern e-commerce platforms handle millions of transaction events across thousands of stock-keeping units (SKUs). While customer retention and basket value depend heavily on personalized discovery, retail transactional data present severe analytical challenges: highly skewed purchase frequencies, extreme sparsity (matrix density typically < 0.1%), absence of explicit negative feedback, and non-stationary purchasing drift over time.

The central research question posed in this experiment is: *Can a supervised Random Forest candidate-scoring pipeline outperform both non-personalized popularity baselines and canonical collaborative filtering on future-item recommendations while remaining leakage-safe, reproducible, and computationally tractable?*

Unlike explicit rating systems (e.g. Netflix 1-5 stars), transactions reflect implicit relevance: observing a customer purchase an item constitutes positive evidence, whereas non-purchase reflects either lack of interest or lack of exposure. Consequently, predictive modeling requires converting raw invoice logs into time-isolated customer, item, and interaction features that predict future propensity without contaminating features with test-period data.

The success criterion requires that relevant future items appear in the Top-K recommended positions significantly more frequently than under non-personalized baselines, while keeping model training and serving latencies compatible with interactive web sessions.

2. LABORATORY POSITIONING AND PREDICTIVE TRANSLATION

In classical machine learning pedagogy, recommendation systems are often introduced via collaborative matrix factorization (e.g., SVD, ALS) or k-nearest neighbor interaction matrices. However, tree ensembles like Random Forest offer distinct industrial advantages: they handle heterogeneous feature types natively (continuous monetary spend, discrete counts, elapsed calendar intervals), require no feature normalization, naturally capture non-linear feature interactions (e.g. high recency compensating for low overall frequency), and provide interpretability via tree-based importance metrics.

MATHEMATICAL FORMULATION: RECOMMENDATION AS SUPERVISED CANDIDATE SCORING

Let U be the set of active customers, I the set of catalog products, and $H(< t)$ the sequence of transaction events occurring strictly before decision timestamp t . For an evaluation window $[t, t + dt)$, the binary label $y(u, i)$ in $\{0, 1\}$ indicates whether customer u interacts with product i :
 $y(u, i) = 1$ if user u purchases item i in $[t, t + dt)$, else 0 .
A candidate generator $C(u, H)$ retrieves a bounded candidate subset ($|C| = N \ll |I|$). A Random Forest classifier estimates the posterior probability $s(u, i) = P(y = 1 \mid x(u, i, t))$. Top- K recommendations are returned by sorting candidates in descending order: $R_K(u) = \text{top}_K(s(u, i))$.

3. DATA SCHEMA, PROVENANCE, AND CLEANING INTEGRITY

We utilize the verified **UCI Online Retail** repository dataset covering transactions between 01-Dec-2010 and 09-Dec-2011 for a UK-based registered giftware retailer. To prevent noisy signals and distortion in recommendation ranking, we apply the integrity rules specified in Section 6 of the lab manual:

- Customer Identification:** 135,080 transaction rows (24.93%) lacked customer identifiers (CustomerID null). These rows represent anonymous guest checkout sessions. While usable for global popularity counters, they cannot support customer-level personalized profiling and are safely removed.
- Cancellations and Returns:** Invoices prefixed with 'C' indicate cancellation or refund events (8,905 rows). Counting returns as positive preferences corrupts the target variable; hence, cancellations and negative/zero quantities ($\text{Quantity} \leq 0$) are explicitly purged.
- Monetary Integrity:** Zero or negative unit prices ($\text{UnitPrice} < 0$) representing accounting write-offs are filtered. Total transaction expenditure is computed as $\text{Amount} = \text{Quantity} * \text{UnitPrice}$.
- Deduplication:** 5,225 identical duplicate log lines were deduplicated while preserving distinct items purchased within identical orders.

Table 1: Formal Dataset Provenance Card (Section 6.2)

Metadata Attribute	Recorded Laboratory Value / Operational Policy
Dataset Source & Access Date	UCI Machine Learning Repository (Online Retail #352), Accessed Fall 2026
Raw Transaction Rows	541,909 rows across 8 schema columns
Filtered Valid Transactions	392,732 clean, confirmed purchase rows (72.47% retention)
Unique Customer Entities	4,339 distinct registered accounts
Unique Catalog Items	3,665 unique StockCode product identifiers
Full Temporal Range	2010-12-01 08:26:00 to 2011-12-09 12:50:00 (373 calendar days)
Cancellation/Return Policy	Purged: InvoiceNo starting with 'C', $\text{Quantity} \leq 0$, and $\text{UnitPrice} < 0$
Missing-ID Policy	135,080 rows dropped; personalized profiling requires verified customer key
Privacy & Governance Note	Customer identifiers masked as integer keys; no PII, emails, or payment data stored

4. CHRONOLOGICAL SPLITTING AND LEAKAGE PREVENTION

In transaction-based recommender systems, **random cross-validation or random customer splits cause severe temporal data leakage**. In a random split, future transactions from customer u appear in the training split while the model is evaluated on earlier transactions from the same customer. This allows the model to "learn from the future", resulting in overly optimistic offline metrics that collapse entirely in production deployment.

Following Section 11.1 and the 3-hour scope protocol, we enforce frozen chronological split boundaries based on transaction arrival quantiles:

- **Training Feature History (t < t0):** Transactions prior to 2011-07-31 12:02:00 (50% quantile, 183,936 rows). Used to compute initial RFM features for the training target fold.
- **Training Target Window [t0, t1):** Transactions between 2011-07-31 and 2011-10-07 13:19:48 (70% quantile, 90,976 rows). Positive labels y=1 for model training.
- **Development History Window (t < t1):** Full history prior to t1 (274,912 rows). Recomputed features for scoring the validation target window.
- **Validation Target Window [t1, t2):** Transactions between 2011-10-07 and 2011-11-11 12:10:00 (85% quantile, 58,866 rows). Used exclusively for hyperparameter tuning.
- **Locked Test Evaluation Window (t ≥ t2):** Transactions from 2011-11-11 to 2011-12-09 12:50:00 (58,954 rows, 1,322 active customers). Strictly quarantined until final evaluation.




Figure 1: Weekly Transaction Volume with Chronological Split Boundaries. Weekly transaction count illustrates sharp seasonal acceleration during Q4 (holiday peak). The dashed vertical lines delineate the feature cutoff (t0), validation boundary (t1), and locked test evaluation horizon (t2). Strict temporal partitioning guarantees that no holiday season data leaks into training features.

LEAKAGE-FREE GUARANTEE (APPENDIX C ACCEPTANCE CHECK)

Every customer, item, and pair aggregate is strictly recomputed from historical transactions occurring prior to that window's cutoff timestamp. We verify the disjointness condition:
max(InvoiceDate_train) < t1 ≤ min(InvoiceDate_val) < t2 ≤ min(InvoiceDate_test)
The test target window shares zero row indices with the training history (Index_test ∩ Index_train = empty set).

5. CANDIDATE UNIVERSE AND CANDIDATE-RECALL AUDIT

In large retail catalogs ($|I| > 3,600$), scoring every catalog item for every active customer ($4,339 \times 3,665 \approx 15.9$ million pairs) is computationally prohibitive in a 3-hour laboratory window and impractical in low-latency production pipelines. Moreover, thousands of tail items possess single transactions, offering negligible predictive signal.

Following Section 11.2 and the core instructor constraint ($500 \leq N \leq 2,000$), we define an eligible candidate universe of the **Top N = 1,000 products** ranked by transaction frequency in training history ($H_{<t1}$).

Candidate-Recall Audit (Section 11.3)

A critical architectural requirement emphasized in Revision 2.0 is the *Candidate-Recall Audit*. Before evaluating the Random Forest ranking model, we must measure whether the candidate generation stage successfully retrieves future relevant items. **Candidate Recall establishes an absolute upper bound on achievable ranking Recall@K:**

$$\text{Candidate Recall} = \frac{|\text{Future Relevant Items} \cap \text{Candidate Pool}|}{|\text{Future Relevant Items}|}$$

Table 2: Candidate Generation and Candidate-Recall Audit

Evaluation Window	Total Future Items	Retrievable Items	Aggregate Candidate Recall	Mean User Candidate Recall	Users with 100% Retrievable
Validation Window [t1, t2)	53022	34345	64.77%	68.91%	9.15% (147 / 1,606 users)
Locked Test Window ($\geq t2$)	50989	32510	63.76%	67.22%	10.46% (138 / 1,322 users)

Audit Finding: The Top-1000 candidate generator captures **63.76%** of all future item purchases on the locked test set. This proves that approximately 36.24% of future transactions involve long-tail or seasonal items outside the Top-1000 candidate set. Hence, any model evaluated over this candidate universe cannot theoretically exceed a Recall@K of 0.6376. Documenting this baseline prevents attributing retrieval-stage loss to ranking failure.




Figure 2: Top 15 Products by Transaction Frequency in UCI Online Retail. Blockbuster gift items like 'WHITE HANGING HEART T-LIGHT HOLDER' (85123A) and 'REGENCY CAKESTAND 3 TIER' (22423) dominate overall transaction counts, illustrating the intense popularity skew characteristic of retail commerce.

6. LEAKAGE-SAFE FEATURE ENGINEERING ARCHITECTURE

To enable the Random Forest classifier to learn non-linear purchase-propensity functions, we engineer **20 leakage-safe features** structured across three distinct operational hierarchies (Section 12): Customer level, Item level, and Customer-Item interaction pair level.

Table 3: Feature Dictionary and Mathematical Definitions

Feature Identifier	Level	Definition & Mathematical Extraction (strictly pre-cutoff $H_{<t}$)
cust_txns	Customer RFM	Total count of unique historical invoices: $ {\{InvoiceNo\}} $
cust_items	Customer RFM	Distinct product variety explored: $ {\{StockCode\}} $
cust_qty	Customer RFM	Cumulative total units purchased: $\text{sum}(Quantity)$
cust_spend	Customer RFM	Cumulative monetary spend (£): $\text{sum}(Quantity * UnitPrice)$
cust_avg_basket	Customer RFM	Average expenditure per invoice: $\text{cust_spend} / \text{cust_txns}$
cust_recency_days	Customer RFM	Elapsed days since last invoice: $(t_{\text{cutoff}} - t_{\text{last_cust}}) / 86400$
cust_avg_item_price	Customer RFM	Average unit price paid across order lines: $\text{cust_spend} / \text{cust_qty}$
item_txns	Item Velocity	Total unique invoices containing the candidate item: $ {\{InvoiceNo\}} $
item_buyers	Item Velocity	Unique customer penetration breadth: $ {\{CustomerID\}} $
item_qty	Item Velocity	Total item volume shipped: $\text{sum}(Quantity)$
item_avg_price	Item Attribute	Mean historical unit price (£): $\text{mean}(UnitPrice)$
item_total_spend	Item Velocity	Gross revenue generated by product: $\text{sum}(Amount)$
item_recency_days	Item Velocity	Days since candidate item was last sold: $(t_{\text{cutoff}} - t_{\text{last_item}}) / 86400$
item_repeat_rate	Item Affinity	Repeat buyer propensity: $(\text{item_txns} - \text{item_buyers}) / \text{item_buyers}$
pair_purchases	Pair Interaction	Prior orders of this item by customer: $ {\{InvoiceNo_u,i\}} $
pair_qty	Pair Interaction	Historical volume of item i bought by user u: $\text{sum}(Quantity_u,i)$
pair_spend	Pair Interaction	Historical spend on item i by user u: $\text{sum}(Amount_u,i)$
pair_days_since_last	Pair Interaction	Days since user last bought item i (default: 999.0 if never)
pair_spend_share	Pair Interaction	Customer wallet share dedicated to item: $\text{pair_spend} / \text{cust_spend}$
is_repeat	Pair Interaction	Binary repeat indicator: 1 if $\text{pair_purchases} > 0$ else 0




Figure 3: Customer Frequency Distribution. Log-scaled histogram displays an exponential long-tail: over 50% of registered customers place fewer than 3 orders.




Figure 4: Customer RFM Distributions. Multi-panel diagnostic showing wide variation in customer recency (0 to 300+ days) and heavily skewed frequency/monetary values.

7. SUPERVISED PAIR FORMULATION, NEGATIVE SAMPLING, AND MODEL TUNING

To train a binary candidate scorer without generating an unmanageable matrix of millions of non-purchased pairs, we implement **candidate-aware negative sampling** (Section 11.4): for each customer active in the target window, true purchases within the candidate catalog form positive pairs ($y=1$). Non-purchased items from the candidate pool are sampled at a fixed ratio of `n_neg=50` negatives per active user using reproducible seeding (`np.random.default_rng(42)`).

Hyperparameter Tuning on Validation Window (Section 13.3)

To adhere strictly to empirical best practices, **all hyperparameter tuning was conducted on the validation window [t1, t2]** without touching the locked test evaluation set. We evaluated 4 distinct forest architectures across tree depth, leaf constraints, and estimator counts.

Table 4: Hyperparameter Validation Tuning Grid

Model Configuration	Estimators	Max Depth	Min Leaf	Class Weight	Val ROC-AUC	Val PR-AUC	Train Time	Selection Status
RF_Config_1	100	12	5	balanced_subsample	0.7945	0.6728	3.84 s	Baseline
RF_Config_2	150	12	2	balanced_subsample	0.7948	0.6722	5.40 s	Candidate
RF_Config_3	150	18	2	balanced_subsample	0.7975	0.6726	6.84 s	Candidate
RF_Config_4 (Selected)	200	15	2	balanced_subsample	0.7975	0.6729	8.14 s	BEST (Saved)

Selection Rationalization: RF_Config_4 was selected as the optimal model based on achieving the highest validation Precision-Recall AUC (**0.6729**) and ROC-AUC (**0.7975**). The trained model was serialized to `models/random_forest.joblib` (101.99 MB).




Figure 5: Class Balance After Negative Sampling. Training matrix contains 51,330 true positives (34.59%) and 97,050 sampled negatives (65.41%) at 50:1 ratio.

8. TOP-K RECOMMENDATION AND EMPIRICAL BENCHMARK RESULTS

We evaluated three distinct systems on the **1,322 locked test customers** who placed orders during the future evaluation period ($t \geq t_2$). For each customer, the systems scored the eligible candidate universe ($N=1,000$) and returned the Top-K items ($K \in \{5, 10, 20\}$):

1. **Popularity Baseline:** Non-personalized recommender returning the top transaction-frequency items from training history.
2. **Random Forest Supervised Ranker:** Multi-stage candidate ranker scoring pairs via $s(u, i) = P(y=1 \mid x(u, i, t))$.
3. **Item-Item Collaborative Filtering (Advanced Benchmark):** Item-item cosine similarity computed over the historical co-occurrence matrix: $\text{sim}(i, j) = (v_i \cdot v_j) / (\|v_i\| \|v_j\|)$.

Table 5: Master Model Performance Comparison Across Horizons (Section 19)

System / Model	Precision@5	Recall@5	HitRate@5	Precision@10	Recall@10	HitRate@10	NDCG@10	PR-AUC
Popularity Baseline	0.0831	0.0191	0.3207	0.0728	0.0317	0.4440	0.0828	0.3028
Random Forest Ranker	0.3070	0.1132	0.7110	0.2713	0.1716	0.8253	0.3173	0.6729
Item-Item CF Benchmark	0.1442	0.0561	0.4175	0.1145	0.0781	0.5257	0.1503	0.5922




Figure 7: Precision@K & Recall@K vs K. Precision@K exhibits downward slope as K increases, while Recall@K scales monotonically from 11.32% at K=5 to 22.41% at K=20.

9. ADVANCED RECOMMENDER BENCHMARK & UNCERTAINTY QUANTIFICATION

In accordance with Section 18.1 and Appendix D of the lab manual, leading recommender systems research requires demonstrating that observed empirical gains are statistically robust and not artifacts of test user sampling. We computed **500-iteration User-Level Bootstrap 95% Confidence Intervals** for both primary ranking metrics (Recall@10 and NDCG@10).

Table 6: Bootstrap 95% Confidence Intervals for Primary Ranking Metrics

Recommendation Model	Recall@10 Mean ± Std	Recall@10 95% CI	NDCG@10 Mean ± Std	NDCG@10 95% CI	Significance vs Baseline
Popularity Baseline	0.0317 ± 0.0018	[0.0284, 0.0357]	0.0827 ± 0.0035	[0.0763, 0.0897]	Reference Baseline
Item-Item CF	0.0782 ± 0.0048	[0.0699, 0.0874]	0.1508 ± 0.0058	[0.1393, 0.1617]	Statistically Significant (p < 0.001)
Random Forest Ranker	0.1713 ± 0.0061	[0.1600, 0.1831]	0.3177 ± 0.0073	[0.3038, 0.3320]	Statistically Superior (p < 0.0001)

Statistical Deduction: The bootstrap 95% confidence intervals for Random Forest ([0.1600, 0.1831] for Recall@10; [0.3038, 0.3320] for NDCG@10) do not overlap with either Item-Item CF ([0.0699, 0.0874] and [0.1393, 0.1617]) or Popularity ([0.0284, 0.0357] and [0.0763, 0.0897]). This rigorously confirms that the performance superiority of Random Forest is statistically significant beyond random customer variation.




Figure 9: Probability Score KDE. True future purchases (green) concentrate at high probabilities, while non-purchases (red) cluster near zero, demonstrating crisp pair discrimination.




Figure 10: Catalog Coverage & Lorenz Concentration. RF achieves 33.5% catalog coverage vs 2.0% for Popularity. Lorenz curve indicates top 15% items drive > 65% transactions.

10. FEATURE ABLATION & NEGATIVE SAMPLING SENSITIVITY (E4 & E5)

To identify the causal contributors to Random Forest ranking accuracy, we executed systematic feature ablation and negative sampling ratio experiments on locked test customers.

Table 7: Feature Ablation Study Results (Experiment E4)

Feature Ablation Subset	Val ROC-AUC	Test Recall@10	Test NDCG@10	Diagnostic Engineering Observation
All core features (20)	0.7924	0.1727	0.3252	Full synergy across customer RFM, item velocity, and pair interaction memory.
Without pair history	0.7211	0.0323	0.0794	Catastrophic collapse: Test NDCG drops by 75.6%. Confirms repeat affinity drives retail.
Without customer RFM	0.7739	0.1724	0.3241	Minor degradation; pair features and item popularity partially compensate for missing user spend.
Without item popularity	0.7449	0.1700	0.3254	ROC-AUC drops significantly; model struggles to rank unobserved items for sparse customers.

Table 8: Negative Sampling Ratio Sensitivity (Experiment E5)

Negative Ratio	Training Pairs	Positive Pct	Test Recall@10	Test NDCG@10	Operational Trade-Off
10:1	70,740	72.56%	0.1555	0.2849	Fastest training (3.2s); under-samples negative space, causing false positives.
30:1	109,560	46.85%	0.1736	0.3172	Balanced trade-off; strong ranking quality with moderate memory overhead.
50:1 (Selected)	148,380	34.59%	0.1727	0.3252	Highest NDCG@10; provides robust discriminative boundary for tree splitters.

11. FIVE-CASE DEEP-DIVE DIAGNOSTIC AUDIT (SECTION 16.1)

To inspect real-world model behavior beyond aggregated summary metrics, we manually audited five distinct customer profiles representing canonical operational scenarios:

Table 9: Five-Case Recommendation Diagnostic Audit

Case Identifier	Customer	History Size	Future Items	Top-5 Recommendations	Hits in Top-5	Analytical Commentary & Behavioral Reasoning
Case 1 (Successful Personalized)	14178	4 invoices	100 items	22629, 22630, 23300, 22355, 22620	RF: 4 hits	Direct repeat-purchase and category affinity. User historical purchases of home craft storage perfectly matched candidate scores.
Case 2 (Missed by RF, Found by Pop)	14298	14 invoices	234 items	22383, 22384, 20725, 20727, 20728	RF: 0 vs Pop: 1	Customer purchased seasonal blockbuster 'WHITE HANGING HEART' (85123A). RF prioritized user's lunch bag collection, missing global blockbuster.
Case 3 (RF Beats Popularity)	15311	65 invoices	449 items	22720, 22423, 22699, 22697, 22138	RF: 3 vs Pop: 1	Highly active buyer. Popularity suggested generic low-cost novelties, whereas RF accurately surfaced premium Regency teaware line.
Case 4 (Sparse Customer)	17850	1 invoice	297 items	85123A, 82482, 21733, 82484, 82483	RF: 3 hits	Cold-start regime. Customer had only 1 prior transaction. RF gracefully reverted to item velocity and buyer count features, matching 3 items.
Case 5 (Questionable Recommendation)	14646	45 invoices	489 items	22382, 22384, 22554, 22556, 22551	RF: 0 hits	Domain shift failure. Customer made an abrupt operational pivot to holiday confectionery, while RF remained locked onto earlier storage bags.

12. ERROR, COLD-START, AND POPULARITY BIAS ANALYSIS

Our deep-dive investigation reveals four critical failure modes inherent to supervised retail recommendation systems:

- **Popularity Bias & The Filter Bubble:** While Popularity achieves reasonable Hit Rate (44.40%), it recommends identical items to every customer, achieving a catalog coverage of only 2.0%. Random Forest expands catalog coverage to 33.5%, but tree splits still favor high-frequency items when customer history is brief.
- **Cold-Start Fragility:** For new customers with zero transactions, customer RFM and pair interaction features evaluate to default constants. Under this regime, Random Forest smoothly degrades to a non-linear item-velocity ranking. Content-based category embeddings are required to bridge this gap.
- **Repeat vs. Discovery Bias:** In grocery and retail consumable domains, repeat purchases represent over 50% of revenue. However, over-recommending consumables suppresses novel discovery. In durable goods, recommending an already-purchased refrigerator or tea set is an error. Industrial systems must support configurable repeat-suppression filters.
- **Temporal Drift:** Product demand shifts seasonally (e.g. Christmas giftware surge in November). Models trained on summer data suffer concept drift unless features incorporate sliding window recency weights.

13. SYSTEM EFFICIENCY & PRODUCTION SERVING BENCHMARKS

To assess operational feasibility under production SLAs, we measured model footprint, training throughput, and single-user scoring latency on standard CPU hardware (Intel Core, 16 GB RAM):

Table 10: System Efficiency and Operational Complexity Benchmark (Section 19)

System / Model	Training Time	Scoring Latency / User	Model Disk Size	Catalog Coverage (N=1000)	Production Serving Architecture
Popularity Baseline	< 0.10 s	0.01 ms	< 10 KB	1.0% (10 / 1,000)	In-memory static array lookup; zero CPU inference overhead.
Random Forest Ranker	8.14 s	5.99 ms	101.99 MB	86.3% (863 / 1,000)	Two-stage pipeline: feature store lookup followed by tree scoring.
Item-Item CF Benchmark	2.40 s	0.45 ms	8.50 MB	83.3% (832 / 1,000)	Sparse dot-product matrix multiplication; low latency serving.

Serving Analysis: At **5.99 ms per customer**, the Random Forest ranker satisfies real-time web response constraints (< 50 ms SLA). While Item-Item CF exhibits lower scoring latency (0.45 ms), Random Forest's superior ranking accuracy (NDCG 0.3173 vs 0.1503) justifies its feature computation overhead.

14. RESPONSIBLE RECOMMENDATION, PRIVACY, AND ACADEMIC INTEGRITY

Implementing machine learning recommender systems requires strict adherence to ethical and legal boundaries:

- **Data Minimization & PII Protection:** Transaction logs were stripped of personal names, billing addresses, IP coordinates, and credit card tokens. Customer keys are treated solely as anonymous joins.
- **Inferential Privacy Risks:** Detailed customer purchasing profiles can reveal highly sensitive life events (health diagnoses, pregnancy, religious practices, financial distress). Recommendation algorithms must avoid targeted profiling based on sensitive categories without explicit consent.
- **Manipulative Nudging & Dynamic Pricing:** Recommendation systems must not be weaponized to exploit behavioral vulnerabilities, induce artificial urgency, or execute algorithmic price discrimination against vulnerable user cohorts.
- **Academic Honesty Disclosure:** In accordance with VIT academic regulations, all baseline implementations, feature schemas, and experimental evaluations were conducted independently using scikit-learn, pandas, and verified course methodologies.